

参考文献：<https://blog.csdn.net/yuanlong122716/article/details/105160545/>  
官方文档：<https://docs.spring.io/spring-kafka/reference/html/>

一、简介

- 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
1. kafka是一个分布式-订阅消息传递系统，使用scala语言开发，基于zookeeper进行协调，多分区、多副本；它的特性是高吞吐、可持久化、可水平扩展、支持流数据处理，它具备三大功能：消息系统、存储系统、流式处理系统。

2. kafka基本概念

3 Producer ：生产者，发送消息的一方

4 Consumer ：消费者，接收消息的一方

5 Broker ：kafka节点，一个节点就是一个kafka server进程

6 Topic ：主题，消息以主题来进行归类

7 Partition ：分区，主题的所有消息分布在不同的区中，每个分区的消息一定是不同的，分区可以分布在不同的broker中

8 Replica ：副本机制，每个分区引入多副本，leader副本和follower副本，leader副本处理读写，follower副本负责同步leader副本的数据，出现故障时，follower副本中重新选举出新的leader副本，进行故障转移

9 PacificA ：kafka采用的一致性协议

3. Kafka特性

高吞吐量、低延迟：kafka每秒可以处理几十万条消息，它的延迟最低只有几毫秒，每个topic可以分多个partition，consumer group 对partition进行consume操作；

可扩展性：kafka集群支持热扩展；

持久性、可靠性：消息被持久化到本地磁盘，并且支持数据备份防止数据丢失；

容错性：允许集群中节点失败（若副本数量为n, 则允许n-1个节点失败）；

高并发：支持数千个客户端同时读写；

支持实时在线处理和离线处理：可以使用Storm这种实时流处理系统对消息进行实时进行处理，同时还可以使用Hadoop这种批处理系统进行离线处理；

4. Kafka 使用场景

日志收集：一个公司可以用Kafka可以收集各种服务的log，通过kafka以统一接口服务的方式开放给各种consumer，例如Hadoop、Hbase、Solr等；

消息系统：解耦和生产者和消费者、缓存消息等；

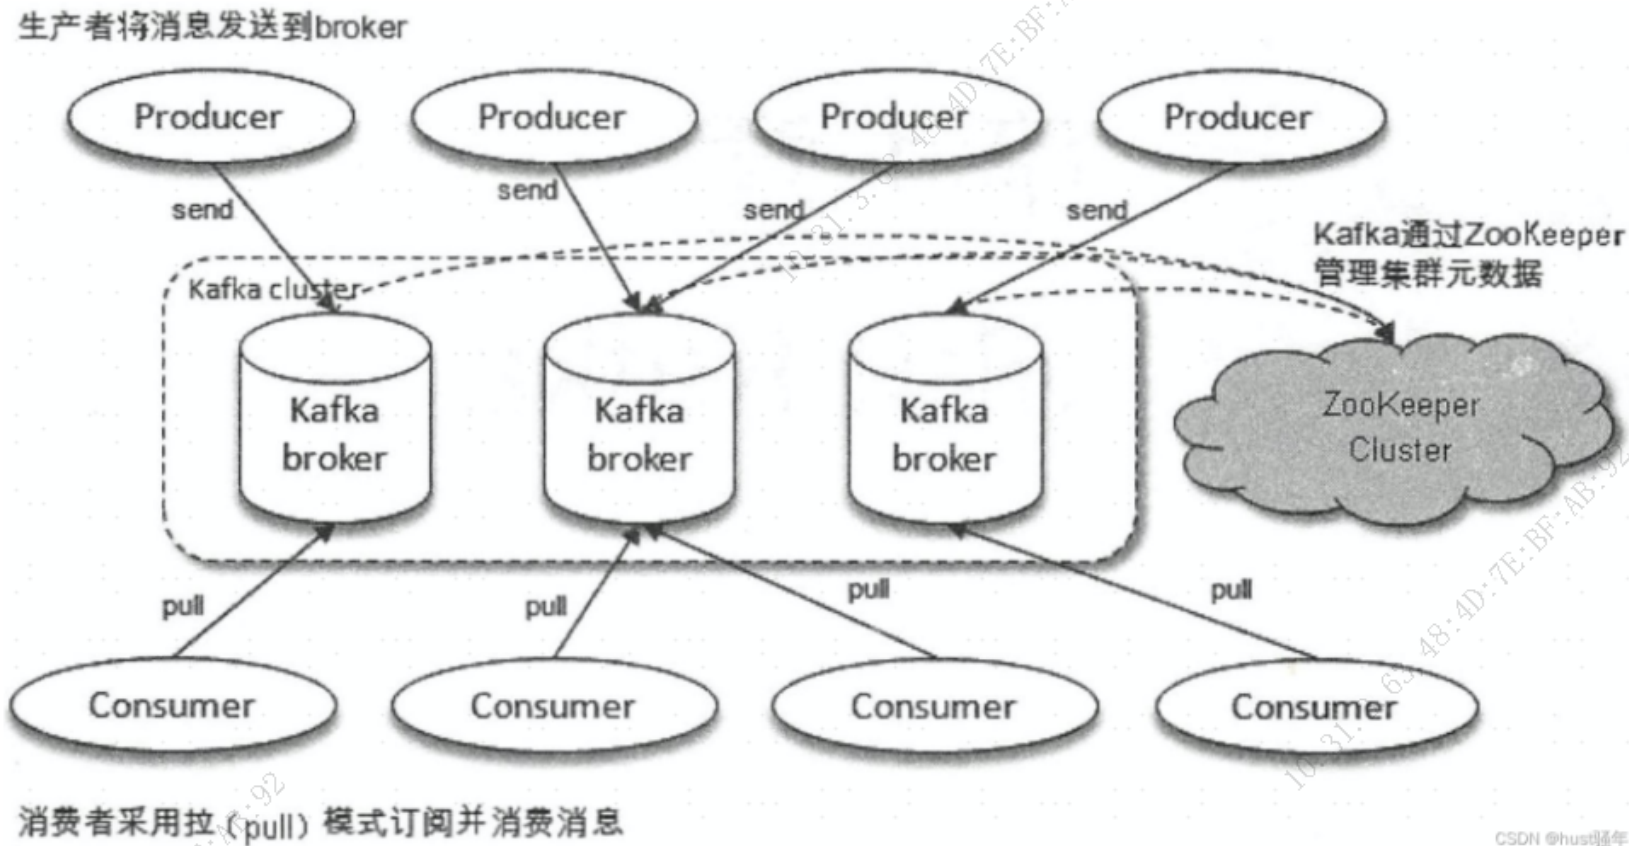
用户活动跟踪：Kafka经常被用来记录web用户或者app用户的各种活动，如浏览网页、搜索、点击等活动，这些活动信息被各个服务器发布到kafka的topic中，然后订阅者通过订阅这些topic来做实时的监控分析，或者离线分析和挖掘；

运营指标：Kafka也经常用来记录运营监控数据。包括收集各种分布式应用的数据，生产各种操作的集中反馈，比如报警和报告；

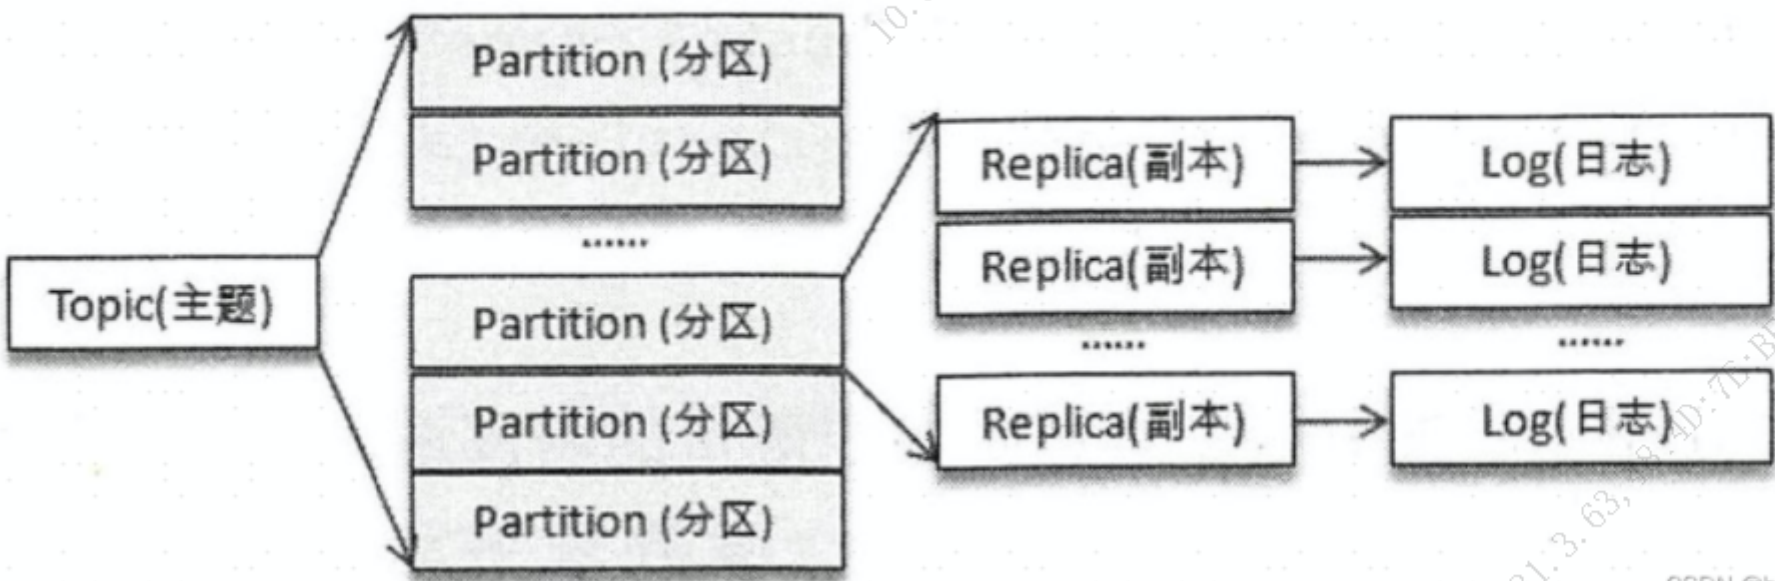
流式处理：比如spark streaming和storm；

事件源：

整体架构



主题与分区



二、环境搭建

```
1 1. 安装zk
2 //docker stop zookeeper; docker rm zookeeper
3 docker run -d --name zookeeper -p 2181:2181 -v /etc/localtime:/etc/localtime wurstmeister/zookeeper
4
5 2. 安装kafka
6 //docker stop kafka; docker rm kafka
7 docker run -d --name kafka \
8 -p 9092:9092 \
9 -e KAFKA_BROKER_ID=0 \
10 -e KAFKA_ZOOKEEPER_CONNECT=10.207.0.167:2181/kafka \
```



```

11 -e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://10.207.0.167:9092 \
12 -e KAFKA_LISTENERS=PLAINTEXT://0.0.0.0:9092 \
13 -t wurstmeister/kafka
14
15 //参数说明
16 -e KAFKA_BROKER_ID=0 在kafka集群中，每个kafka都有一个BROKER_ID来区分自己
17 -e KAFKA_ZOOKEEPER_CONNECT=10.207.0.167:2181/kafka 配置zookeeper管理kafka的路径10.207.0.167:2181/kafka
18 -e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://10.207.0.167:9092 把kafka的地址端口注册给zookeeper
19 -e KAFKA_LISTENERS=PLAINTEXT://0.0.0.0:9092 配置kafka的监听端口
20 -v /etc/localtime:/etc/localtime 容器时间同步虚拟机的时间

```

### 三、环境验证

```

1 1.zk验证
2 //进入容器验证
3 docker exec -it zookeeper bash
4 > cd /opt/zookeeper-3.4.13/bin
5 > ./zkCli.sh //进入zk客户端
6 > ls / //查看根目录有俩节点 [kafka, zookeeper]
7 > ls /kafka/brokers/topics/zero/partitions //查看zero主题的partitions信息
8 > get /kafka/brokers/topics/zero //显示该节点的数据内容和属性信息
9 > ls2 /kafka/brokers/topics/zero //显示该节点的子节点信息和属性信息

```

```
2022-06-07 07:07:18,285 [myid:] - INFO [main-Environment@60] - Client-Environment: user.dir=/opt/zookeeper-3.4.13/bin
2022-06-07 07:07:18,284 [myid:] - INFO [main:ZooKeeper@442] - Initiating client connection, connectString=localhost:2181 sessionTimeout=30000 watcher=org.apache.zookeeper.ZooKeeperMain$MyWatcher@7ab9aala
Welcome to ZooKeeper!
2022-06-07 07:07:18,285 [myid:] - INFO [main-SendThread(localhost:2181):ClientCnxn$SendThread@1029] - Opening socket connection to server localhost/127.0.0.1:2181. Will not attempt to authenticate using SASL (unknown error)
JLine support is enabled
2022-06-07 07:07:18,294 [myid:] - INFO [main-SendThread(localhost:2181):ClientCnxn$SendThread@879] - Socket connection established to localhost/127.0.0.1:2181, initiating session
[zookeeper:localhost:2181(CONNECTING) 0] 2022-06-07 07:07:18,319 [myid:] - INFO [main-SendThread(localhost:2181):ClientCnxn$SendThread@1303] - Session establishment complete on server localhost/127.0.0.1:2181, sessionId = 0x10004a
fb6130003, negotiated timeout = 30000

WATCHER:

WatchedEvent state:SyncConnected type:None path:null
ls /
[kafka, zookeeper]
[zookeeper:localhost:2181(CONNECTED) 1]
```

```

1 2.kafka验证
2 //进入容器验证
3 docker exec -it kafka bash
4 > cd /opt/kafka_2.13-2.8.1/bin
5 > ./kafka-console-producer.sh --broker-list localhost:9092 --topic zero //生产者指定topic发送消息
6 > {"code": 200, "message": "success", "data":{"list":[1,2,3]}}
7 > ./kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic zero --from-beginning //另外打开窗口监听(消费)
8 > ./kafka-topics.sh --create --zookeeper 10.207.0.167:2181/kafka --topic tp-zero1 --replication-factor 1 --partitions 2 //创建topic和partitions
9 > ./kafka-topics.sh --delete --zookeeper 10.207.0.167:2181/kafka --topic tp-zero1 //同时删除容器中的topic数据和zk的topic目录
10 > cat /opt/kafka_2.13-2.8.1/config/server.properties //查看配置信息，如日志目录为 /kafka/kafka-logs-4432dfb12cf7
11 > cd /kafka/kafka-logs-4432dfb12cf7/zero-0; cat 00000000000000000000.log

```

zk

```
WatchedEvent state:SyncConnected type:None path:null
[zk: localhost:2181(CONNECTED) 0] ls /kafka/brokers/topics
[tp-zero1, consumer_offsets, zero, sun]
[zk: localhost:2181(CONNECTED) 1] ls /kafka/brokers/topics
[_consumer_offsets, zero, sun]
[zk: localhost:2181(CONNECTED) 2]
```

kafka

```
root@4432dfb12cf7: /# cd /kafka/kafka-logs-4432dfb12cf7
root@4432dfb12cf7: /kafka/kafka-logs-4432dfb12cf7# ls
__consumer_offsets-0  __consumer_offsets-15  __consumer_offsets-21  __consumer_offsets-28  __consumer_offsets-34  __consumer_offsets-40  __consumer_offsets-47  __consumer_offsets-9  zero-0
__consumer_offsets-1  __consumer_offsets-16  __consumer_offsets-22  __consumer_offsets-29  __consumer_offsets-35  __consumer_offsets-41  __consumer_offsets-48  cleaner-offset-checkpoint
__consumer_offsets-10  __consumer_offsets-17  __consumer_offsets-23  __consumer_offsets-3  __consumer_offsets-36  __consumer_offsets-42  __consumer_offsets-49  log-start-offset-checkpoint
__consumer_offsets-11  __consumer_offsets-18  __consumer_offsets-24  __consumer_offsets-30  __consumer_offsets-37  __consumer_offsets-43  __consumer_offsets-5  meta.properties
__consumer_offsets-12  __consumer_offsets-19  __consumer_offsets-25  __consumer_offsets-31  __consumer_offsets-38  __consumer_offsets-44  __consumer_offsets-6  recovery-point-offset-checkpoint
__consumer_offsets-13  __consumer_offsets-2  __consumer_offsets-26  __consumer_offsets-32  __consumer_offsets-39  __consumer_offsets-45  __consumer_offsets-7  replication-offset-checkpoint
__consumer_offsets-14  __consumer_offsets-20  __consumer_offsets-27  __consumer_offsets-33  __consumer_offsets-46  __consumer_offsets-8  sun-0
root@4432dfb12cf7: /kafka/kafka-logs-4432dfb12cf7# cd zero-0/
root@4432dfb12cf7: /kafka/kafka-logs-4432dfb12cf7/zero-0# ls
00000000000000000000.index  00000000000000000000000000000000.log  00000000000000000000000000000000.timeindex  leader-epoch-checkpoint
root@4432dfb12cf7: /kafka/kafka-logs-4432dfb12cf7/zero-0# cat 00000000000000000000000000000000.log
{"code": 200, "message": "success", "data": {"list": [1, 2, 3]}}
{"code": 200, "message": "success", "data": {"list": [1, 2, 3]}}
```

#### 四、Springboot整合Kafka

```
1 1.依赖
2 <!-- parent -->
3 <parent>
```


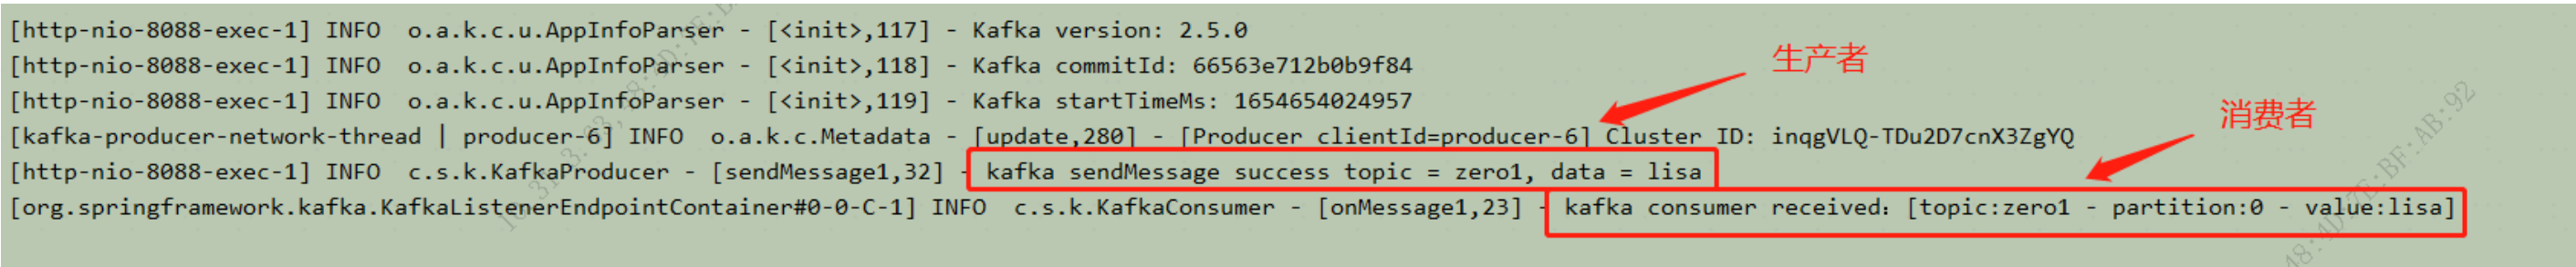


```
4     <groupId>org.springframework.boot</groupId>
5     <artifactId>spring-boot-starter-parent</artifactId>
6     <version>2.3.0.RELEASE</version>
7 </parent>
8 <!-- kafka -->
9 <dependency>
10     <groupId>org.springframework.kafka</groupId>
11     <artifactId>spring-kafka</artifactId>
12 </dependency>
13 <!-- data 同步事务必须要有该依赖 -->
14 <dependency>
15     <groupId>org.springframework.data</groupId>
16     <artifactId>spring-data-commons</artifactId>
17 </dependency>
18 <!-- fastjson -->
19 <dependency>
20     <groupId>com.alibaba</groupId>
21     <artifactId>fastjson</artifactId>
22     <version>1.2.79</version>
23 </dependency>
24
25 2.yml配置
26 spring:
27     kafka:
28         # 集群信息，多个用逗号隔开
29         bootstrap-servers: 10.207.0.167:9092
30         # 【生产者】
31         producer:
32             # 事务ID=事务前缀+1，不为空即开启事务，生效必须  retries>0, acks=all
33             #transaction-id-prefix: kfk_tx
34             # 重试次数
35             retries: 3
36             # 应答级别:多少个分区副本备份完成时向生产者发送ack确认(可选0、1、all/-1)
37             #procedure要求leader在考虑完成请求之前收到的确认数，用于控制发送记录在服务端的持久化，其值可以为如下：
38             #acks = 0 如果设置为零，则生产者将不会等待来自服务器的任何确认，该记录将立即添加到套接字缓冲区并视为已发送。在这种情况下，无法保证服务器已收到记录，并且重试配置将不会生效（因为客户端通常不：
39             #acks = 1 这意味着leader会将记录写入其本地日志，但无需等待所有副本服务器的完全确认即可做出回应，在这种情况下，如果leader在确认记录后立即失败，但在将数据复制到所有的副本服务器之前，则记录将
40             #acks = all 这意味着leader将等待完整的同步副本集以确认记录，这保证了只要至少一个同步副本服务器仍然存活，记录就不会丢失，这是最强有力的保证，这相当于acks = -1的设置。
41             acks: 1
42             # 每次批量发送消息的数量，16K
43             batch-size: 16384
44             # linger.ms为0表示每接收到一条消息就提交给kafka,这时候batch-size其实就没用了
45             properties:
46                 linger.ms: 0
47                 # 自定义分区器，实现Partitioner接口
48                 # partitioner:
49                     # class: com.zeor.producer.CustomizePartitioner
50             # 生产端缓冲区大小，32M
51             buffer-memory: 33554432
52             # 指定消息key和消息体的编解码方式，也可自定义
53             key-serializer: org.apache.kafka.common.serialization.StringSerializer
54             value-serializer: com.sf.kafka.serialization.ObjectSerializer
55         # 【生产者】
56         consumer:
57             # 指定默认消费者group id
58             group-id: zeroGroupId
59             # 是否自动提交offset
60             enable-auto-commit: false
61             properties:
62                 # 消费会话超时时间(超过这个时间consumer没有发送心跳,就会触发rebalance操作)
63                 session.timeout.ms: 120000
64                 # 消费请求超时时间
65                 request.timeout.ms: 180000
66             # 默认为500，批量消费每次最多消费多少条消息
67             #max-poll-records: 50
68             # 提交offset延时(接收到消息后多久提交offset)
69             auto-commit-interval: 100
70             # 当kafka中没有初始offset或offset超出范围时将自动重置offset
71             # earliest:当各分区下有已提交的offset时，从提交的offset开始消费；无提交的offset时，从头开始消费，避免消息丢失
72             # latest:重置为分区中最新的offset(消费分区中新产生的数据)；
73             # none:只要有一个分区不存在已提交的offset,就抛出异常；
74             auto-offset-reset: earliest
75             # 指定消息key和消息体的编解码方式
76             key-deserializer: org.apache.kafka.common.serialization.StringDeserializer
77             value-deserializer: com.sf.kafka.serialization.ObjectDeserializer
78         listener:
79             # 指定listener容器中的线程数（同时消费的监听器），用于提高并发量，建议与分区数量一致
80             concurrency: 3
81             # 消费端监听的topic不存在时，项目启动会报错(关掉)
82             missing-topics-fatal: false
83             # 设置批量消费
84             #type: batch
85             # ACK模式: batch、record、time、count、count_time、manual、manual_immediate
86             # 默认batch，手动调用Acknowledgment.acknowledge()后立即提交，一般使用这种
87             ack-mode: manual_immediate
88
```

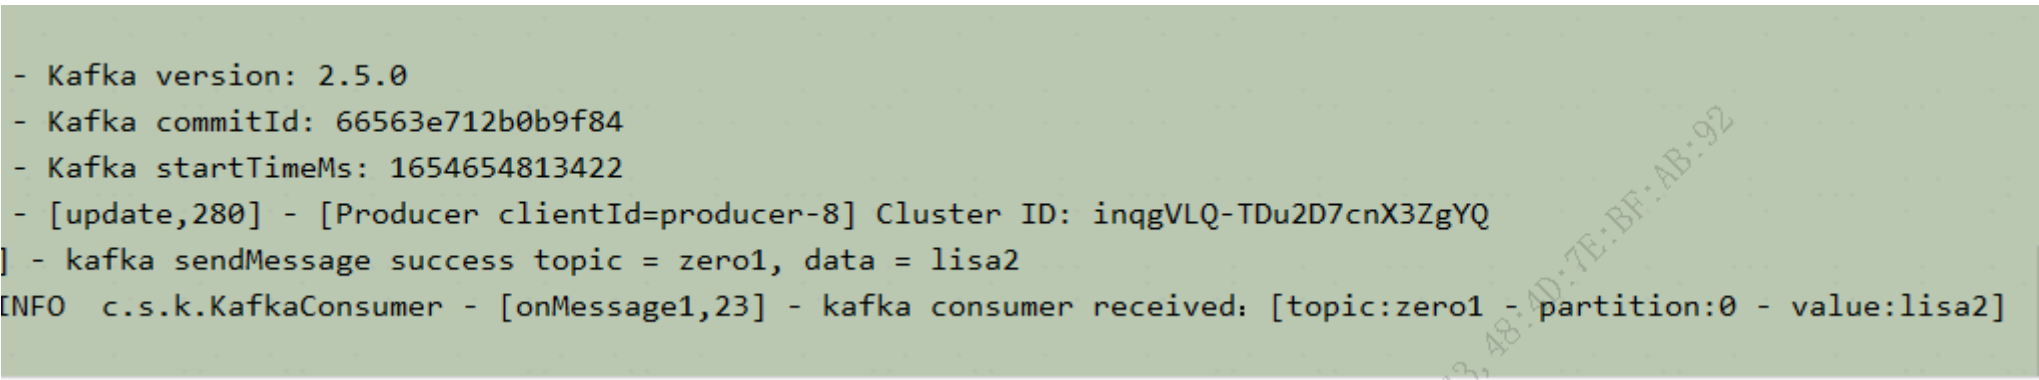
```
89 # 自定义配置
90 kafka:
91   topic:
92     group-id: zeroGroupId
93     topic-name:
94       - zero
95       - zero1
96       - sun
```

五、测试

```
1 1、简单发布-订阅
2 @GetMapping("send/{message}")
3 public void sendMessage1(@PathVariable("message") String message) throws Exception {
4     kafkaTemplate.send(KFK_TOPIC_ZERO1, message).get(10, TimeUnit.SECONDS);
5     log.info("kafka sendMessage success topic = {}, data = {}", KFK_TOPIC_ZERO1, message);
6 }
7
8 //kafkaTopicName = {"zero","zero1","sun"}  topicGroupId = "zeroGroupId"
9 @KafkaListener(topics = "#{kafkaTopicName}", groupId = "#{topicGroupId}")
10 public void onMessage1(ConsumerRecord<String, Object> record) {
11     log.info("kafka consumer received: [topic:{} - partition:{} - value:{}]", record.topic(), record.partition(), record.value());
12 }
```



```
1 2、消费者手动ACK
2 //spring.kafka.consumer.enable-auto-commit=false
3 //spring.kafka.listener.ack-mode=manual_immediate
4 @KafkaListener(topics = "#{kafkaTopicName}", groupId = "#{topicGroupId}")
5 public void onMessage1(ConsumerRecord<String, Object> record, Acknowledgment ack) {
6     log.info("kafka consumer received: [topic:{} - partition:{} - value:{}]", record.topic(), record.partition(), record.value());
7     //手动提交offset
8     ack.acknowledge();
9 }
```



```
1 3、生产者-回调异步发送
2 @GetMapping("send2/{message}")
3 public void sendMessage2(@PathVariable("message") String message) {
4     kafkaTemplate.send(KFK_TOPIC_ZERO1, message).addCallback(success -> {
5         // 消息发送到的topic
6         String topic = success.getRecordMetadata().topic();
7         // 消息发送到的分区
8         int partition = success.getRecordMetadata().partition();
9         // 消息在分区内的offset
10        long offset = success.getRecordMetadata().offset();
11        log.info("发送消息成功:" + topic + "-" + partition + "-" + offset);
12    }, failure -> {
13        log.info("发送消息失败:" + failure.getMessage());
14    });
15 }
16 或
17 @GetMapping("send3/{message}")
18 public void sendMessage3(@PathVariable("message") String message) {
19     ListenableFuture<SendResult<String, Object>> future = kafkaTemplate.send(KFK_TOPIC_ZERO, message);
20     future.addCallback(new ListenableFutureCallback<SendResult<String, Object>>() {
21         @Override
22         public void onFailure(Throwable ex) {
23             log.error("发送消息失败, ex = {}, topic = {}, data = {}", ex, KFK_TOPIC_ZERO, message);
24         }
25     });
26 }
```



```
25 @Override
26     public void onSuccess(SendResult<String, Object> result) {
27         log.info("发送消息成功, topic = {}, data = {}", result.getRecordMetadata().topic(), message);
28     }
29 }
30 }
```

```
- [lambda$sendMessage2$0,45] - 发送消息成功: zero1-0-22
c.s.k.KafkaConsumer - [onMessage1,23] - kafka consumer received: [topic:zero1 - partition:0 - value:lisa11]
```

```
1 4.生产者-事务
2 @GetMapping("send4/{message}")
3 public void sendMessage4(@PathVariable("message") String message){
4     // 不声明事务：后面报错但前面消息已经发送成功了
5     kafkaTemplate.send(KFK_TOPIC_ZERO1, message);
6     throw new RuntimeException("fail");
7 }
8
9 //spring.kafka.producer.transaction-id-prefix=kfk_tx_
10 //spring.kafka.producer.acks=all
11 @GetMapping("send4/{message}")
12 public void sendMessage4(@PathVariable("message") String message){
13     // 声明事务：后面报错消息不会发出去
14     kafkaTemplate.executeInTransaction(operations -> {
15         operations.send(KFK_TOPIC_ZERO1, message);
16         throw new RuntimeException("fail");
17     });
18     log.info("异常消息，发送成功");
19 }
```

//未加事务

```
[Producer clientId=producer-9] Cluster ID: inqgVLQ-TDu2D7cnX3ZgYQ
Consumer - [onMessage1,23] - kafka consumer received: [topic:zero1 - partition:0 - value:会发出去吗]
et.service() for servlet [dispatcherServlet] in context with path [] threw exception [Request proces

异常，依然接收到了

.java:190)
erMethod.java:138)
Handle(ServletInvocableHandlerMethod.java:105)
lerMethod(RequestMappingHandlerAdapter.java:879)
onal(RequestMappingHandlerAdapter.java:793)
```

```
lisa2
lisall
lisall
会发出去吗
```

//加事务

```
o.s.k.s.LoggingProducerListener - [error,254] - Exception thrown when sending a message with key='null' and payload='加事务的消息会发出去吗' to topic zero1:
ce transaction was aborted
TransactionalRequest(Sender.java:422)
.java:312)
239) <1 internal line>
ervlet] - [log,175] - Servlet.service() for servlet [dispatcherServlet] in context with path [] threw exception [Request processing failed; nested exception is ja
a:76)
afkaTemplate.java:463)
internal lines>
roke(InvocableHandlerMethod.java:190)
erRequest(InvocableHandlerMethod.java:138)
```

```
lisa2
lisall
lisall
会发出去吗
加事务的消息
```

```
1 5.同步事务(spring + kafka)
2 @Transactional(transactionManager = "chainedKafkaTransactionManager", rollbackFor = Exception.class)
3 @GetMapping("send5/{message}")
4 public void sendMessage5(@PathVariable("message") String message){
5     //db
6     userService.addUser(new UserBean().setName("丽莎3").setSex("女"));
7     //kafka
8     kafkaTemplate.executeInTransaction(operations -> {
9         operations.send(KFK_TOPIC_ZERO1, message);
10        throw new RuntimeException("fail");
11    });
12    log.info("事务消息，发送成功");
13 }
```

```
Creating a new SqlSession
SqlSession [org.apache.ibatis.session.defaults.DefaultSqlSession@4f76f133] was not registered for synchronization because synchronization is not active
JDBC Connection [com.mysql.cj.jdbc.ConnectionImpl@77d35ca0] will be managed by Spring
==> Preparing: INSERT INTO t_user ( name, sex ) VALUES ( ?, ? )
==> Parameters: 丽莎3(String), 女(String)
<== Updates: 1
Closing non transactional SqlSession [org.apache.ibatis.session.defaults.DefaultSqlSession@4f76f133]
[] 14:38:55.064 [http-nio-8088-exec-2] INFO o.a.k.c.p.ProducerConfig - [logAll,347] - ProducerConfig values:
acks = -1
batch.size = 16384
```

```
producer-kfk_tx_1] INFO o.a.k.c.p.i.TransactionManager - [handleResponse,1509] - [Producer clientId=producer-kfk_tx_1, transactionalId=kfk_tx_1] Discovered transaction coordinator 10.207.
producer-kfk_tx_1] INFO o.a.k.c.p.i.TransactionManager - [setProducerIdAndEpoch,515] - [Producer clientId=producer-kfk_tx_1, transactionalId=kfk_tx_1] ProducerId set to 3 with epoch 3
producer-kfk_tx_1] ERROR o.s.k.s.LoggingProducerListener - [error,254] - Exception thrown when sending a message with key='null' and payload='同步事务test' to topic zero1:
point : Failing batch since transaction was aborted
Sender.maybeSendAndPollTransactionalRequest(Sender.java:422)
Sender.runOnce(Sender.java:312)
Sender.run(Sender.java:239) <1 internal line>
.c.C.[].[.[dispatcherServlet] - [log,175] - Servlet.service() for servlet [dispatcherServlet] in context with path [] threw exception [Request processing failed; nested exception is java
ge5$3(KafkaProducer.java:100)
e.executeInTransaction(KafkaTemplate.java:463)
kaProducer.java:98)
ngCGLIB$$69429ee2.invoke(<generated>)
```

错误回滚

```
lisall
lisall
会发出去吗
加事务的消息
命令输入
```

没用消费到

6	6	小丽	10	女	18642549006	2022-06-01 10:52:09
7	7	小云	20	女	18642549007	2022-06-01 16:52:09
8	8	小静	27	女	18642549006	2022-06-01 16:52:09
9	9	丽莎1	<null>	女	<null>	2022-06-08 11:03:53
10	10	丽莎1	<null>	女	<null>	2022-06-08 11:08:58
11	11	丽莎2	<null>	女	<null>	2022-06-08 11:10:14

丽莎3没入库，同步事务成功

//注释 throw new RuntimeException("fail")

```
lisall
lisall
会发出去吗
加事务的消息
同步事务test
```

7	7	小云	20	女	18642549007	2022-06-01 16:52:09
8	8	小静	27	女	18642549006	2022-06-01 16:52:09
9	9	丽莎1	<null>	女	<null>	2022-06-08 11:03:53
10	10	丽莎1	<null>	女	<null>	2022-06-08 11:08:58
11	11	丽莎2	<null>	女	<null>	2022-06-08 11:10:14
12	15	丽莎3	<null>	女	<null>	2022-06-08 14:44:39

```
1 6.指定颗粒度消费
2 //指定topic、partition、offset消费
3 //同时监听topic1和topic2，监听topic1的0号分区、topic2的 "0号和1号" 分区，指向1号分区的offset初始值为8
4 @KafkaListener(id = "", groupId = "zeroGroupId", topicPartitions = {
5     @TopicPartition(topic = "zero", partitions = {"0"}),
6     @TopicPartition(
7         topic = "zero1",
8         partitions = "0",
9         partitionOffsets = @PartitionOffset(partition = "1", initialOffset = "8")
10    )
11 })
12 public void onMessage2(ConsumerRecord<?, ?> record) {
13     log.info("topic: {} | partition:{} | offset:{} | value:{}", record.topic(), record.partition(), record.offset(), record.value());
14 }
```

```
INFO o.a.k.c.p.i.TransactionManager - [setProducerIdAndEpoch,515] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0]
sendMessage1,41] - kafka sendMessage success topic = zero1, data = aaa
ainer#0-1-C-1] INFO c.s.k.KafkaConsumer - [onMessage2,39] - topic: zero1 | partition:0 | offset:40 | value:aaa
```

```
1 7.发送list数据
2 //spring.kafka.producer.value-serializer=com.sf.kafka.serialization.ObjectSerializer
3 //spring.kafka.consumer.value-serializer=com.sf.kafka.serialization.ObjectSerializer
4
5 // 批量发送
6 @Transactional
7 @GetMapping("send6/list")
8 public void sendMessage6() throws Exception {
9     List<UserBean> list = Arrays.asList(new UserBean().setName("路西"), new UserBean().setName("娜美"));
10    kafkaTemplate.send(KFK_TOPIC_ZERO1, list).get(2, TimeUnit.SECONDS);
11    log.info("kafka sendMessage success topic = {}, data size = {}", KFK_TOPIC_ZERO1, list.size());
12 }
13
14 //List来接收
15 @KafkaListener(groupId = "zeroGroupId", topics = "zero1")
16 public void onMessage3(List<UserBean> list) {
17     for (UserBean record : list) {
```



```
18         log.info(record.toString());
19     }
20 }
```

```
tadata - [update,280] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0] Cluster ID: inqgVLQ-TDu2D7cnX3ZgYQ
i.TransactionManager - [handleResponse,1509] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0] Discovered transaction coordi
i.TransactionManager - [setProducerIdAndEpoch,515] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0] ProducerId set to 2 wi
] - kafka sendMessage success topic = zero1, data size = 2
NFO   c.s.k.KafkaConsumer - [onMessage3,47] - UserBean(id=null, name=路西, age=null, sex=null, phone=null, createDate=null)
NFO   c.s.k.KafkaConsumer - [onMessage3,47] - UserBean(id=null, name=娜美, age=null, sex=null, phone=null, createDate=null)
```

```
1  8.异常降级
2  //异常处理器
3  @Bean
4  public ConsumerAwareListenerErrorHandler consumerAwareErrorHandler() {
5      return (message, exception, consumer) -> {
6          log.info("自定义异常处理器, 处理异常信息: " + message.getPayload());
7          return null;
8      };
9  }
10
11 //指定异常处理
12 @KafkaListener(topics = "zero1", errorHandler = "consumerAwareErrorHandler")
13 public void onMessage4(List<ConsumerRecord<?, ?>> records) {
14     log.info(">>>批量消费一次, records.size()=" + records.size());
15     for (ConsumerRecord<?, ?> record : records) {
16         log.info(record.value().toString());
17     }
18 }
```

```
c.p.i.TransactionManager - [handleResponse,1509] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0] Discovered transaction coordinator 10.207.0.167:9092 (id: 0 r
c.p.i.TransactionManager - [setProducerIdAndEpoch,515] - [Producer clientId=producer-kfk_tx_0, transactionalId=kfk_tx_0] ProducerId set to 2 with epoch 23
,112] - kafka sendMessage success topic = zero1, data size = 2
1] INFO   c.s.k.KafkaConsumer - [onMessage4,54] - >>>批量消费一次, records.size()=2
1] INFO   c.s.k.c.KafkaTopicConfiguration - [lambda$consumerAwareErrorHandler$0,40] - 自定义异常处理器, 处理异常信息: [UserBean(id=null, name=路西, age=null, sex=null, phone=null, c
```

【异常】

- 问题: org.apache.kafka.clients.consumer.ConsumerRecord; nested exception is org.springframework.core.convert.ConversionFailedException: Failed to convert from type [java.util.ArrayList<?>] to type[org.apache.kafka.clients.consumer.ConsumerRecord<?, ?>] for value '[789]';

解决: 关闭批量消费, 注释掉 spring.kafka.consumer.listener.type=batch
- 问题: Producer factory does not support transactions | Must set acks to all in order to use the idempotent producer. Otherwise we cannot guarantee idempotence.

解决: 开启事务, spring.kafka.producer.transaction-id-prefix=kfk\_tx    spring.kafka.producer.acks=all
- 问题: Caused by: java.lang.ClassNotFoundException: org.springframework.data.transaction.ChainedTransactionManager

解决: 添加依赖, spring-data-commons
- 问题: No transaction is in process; possible solutions: run the template operation within the scope of a template.executeInTransaction() operation, start a transaction with @Transactional before invoking the template method

解决: 开启事务后, 要么使用 kafkaTemplate.executeInTransaction(todo...) 执行, 要么在方法上添加事务注解 @Transactional
- 问题: java.lang.ClassCastException: java.util.Arrays\$ArrayList cannot be cast to java.lang.String

解决: 配置的序列化器为StringSerializer, 无法将List转String。自定义Object序列化器, 并在配置里指定